Microsoft Dynamics 365

Dynamics 365 FastTrack Architecture Insights

Performance Best Practices for finance and operations apps

February 15, 2023

Haytham Said
Rich Black

Agenda

- The Performance Testing Approach
- Tools
- Best practices
- Q&A

The Performance Testing Approach

Why

- Ensure Scalability

What

- Throughput
- Overlapping Workload
- Response Times
- Cutover
- Peak Hour

Where

- Pre-Production Scalability
- Production Load Validation

When

- Early
- Late
- With Time to Tune…
- Ongoing

Performance is not our responsibility; it is a cloud solution.

- A customer who will likely move the Go-Live date

Performance Contributions

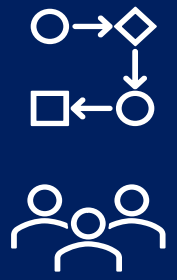

Customer

- Processes and Operations
- Users
- Analytics and Reporting
- Integrations
- Customizations
- Data
- Devices
- Independent Software Vendors (ISV)

Microsoft

- Azure Infrastructure
  - Database
  - Services
  - Network
- Standard Solution

We have tested performance of the integration, this is sufficient. We don't need day-in-the-life tests.

A customer who will likely face performance issues with realistic workloads post Go-Live

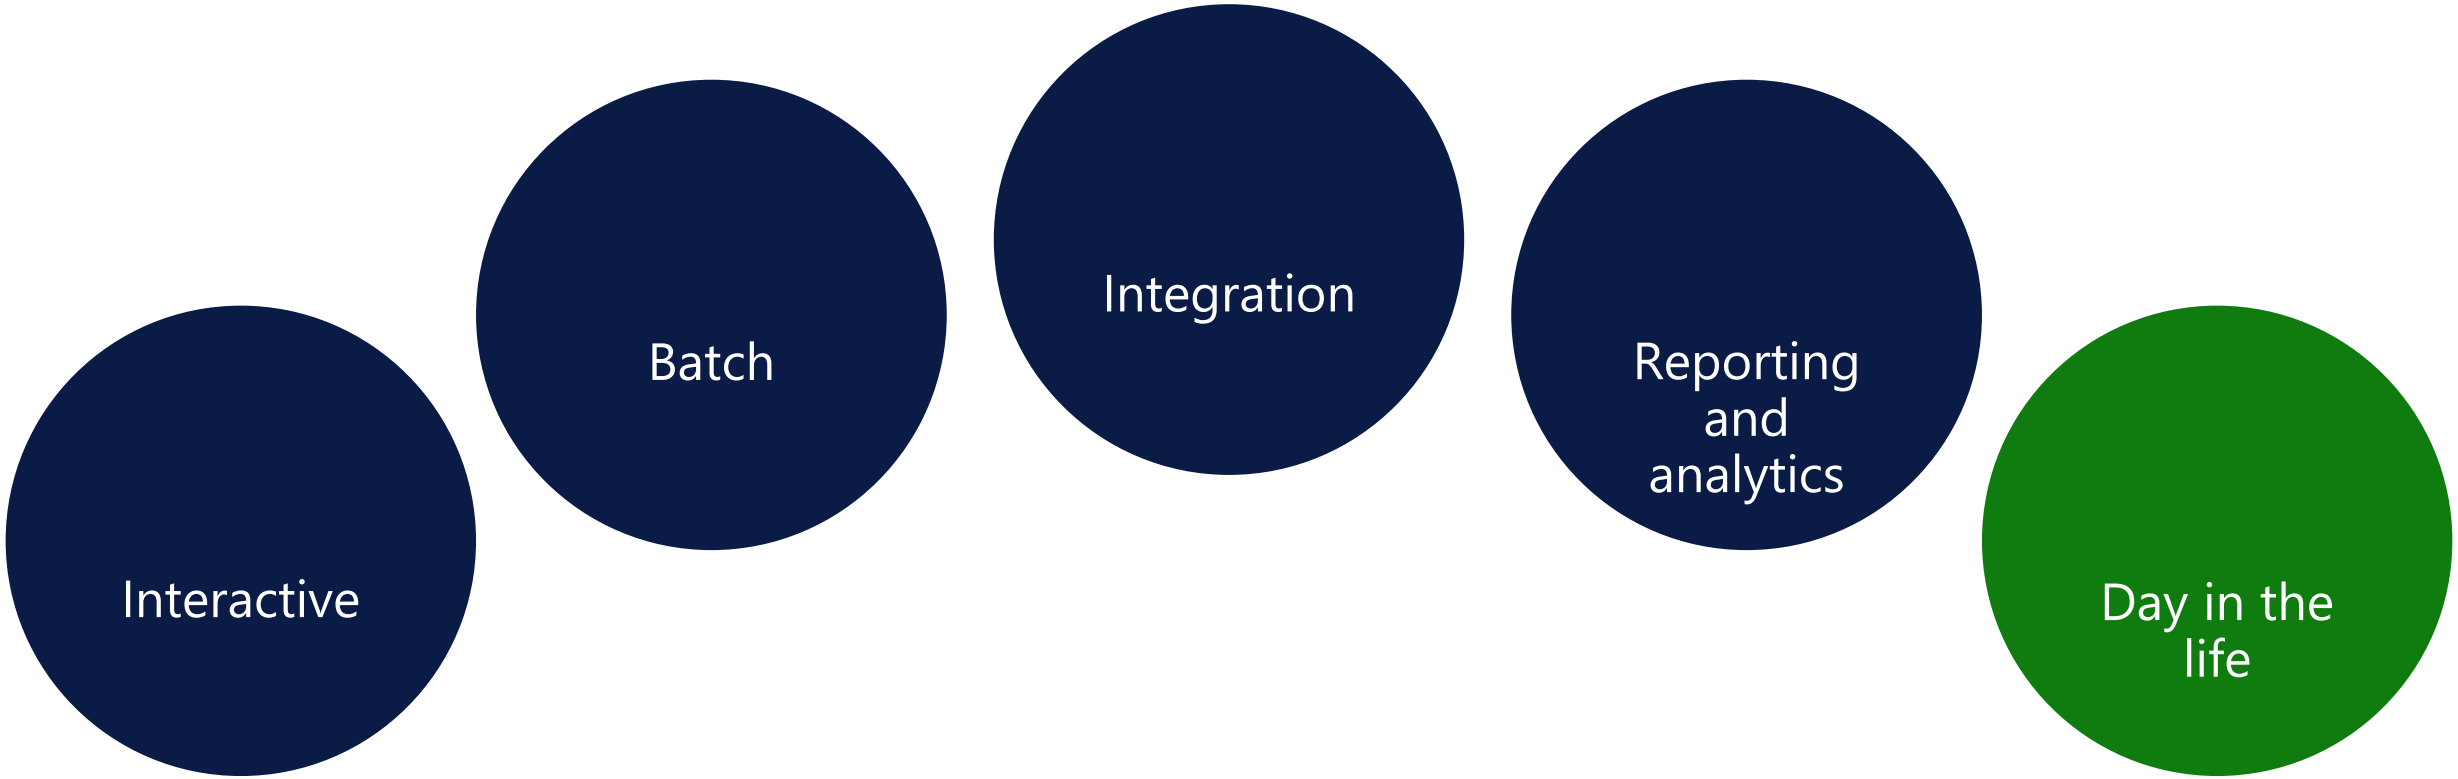

The common types of scenarios

DAY IN THE LIFE SCENARIO

Simulating tasks/activities as they would happen during working hours post Go-Live

Activities like it would happen post Go-Live in PRODUCTION

One unique task tested in isolation works, but testing shouldn't stop there. What if all tasks are initiated at the same time?

We know that tasks and processes will be executed concurrently post Go-Live. If you only test those tasks and processes in isolation, then you are ensuring that you will only discover issues with concurrent task execution in production. (This is not a Good Idea.) It is important to introduce concurrency into your testing!

(i.e. test the interactive, batch, integration and reporting scenarios at the same time, like it would happen during peak hours)

Performance testing is important

- Make it is part of the project plan; this underlines its importance and relevance for the success of the project
- There are multiple dependencies, like resources (environments, people, …), **Go-live criteria**
- It is critically important to define clear goals and non-functional requirements
- Volumes and their mix on a timeline is key

Performance awareness is key

- Call for performance feedback
- Train users to give valid feedback
- Have NFRs to validate against
- Make work items part of the backlog and iterations and define acceptance criteria
- Know the critical areas to allow proper prioritization

Plan iterations

- Performance testing is not a one-step exercise
- There will be findings and questions raised, which need (planned) time, fixes and mitigations need validation and thus, further test executions
- Leverage existing performance metrics (e.g., review as exit criteria)
- Consider one-time activities, like month-end invoicing, year-end activities, and create test scenarios for these, too
- Make sure your sizing estimate is regularly updated and as accurate as possible (after Go-Live)

We will use the PRODUCTION environment for performance testing / stress testing.

A customer who will likely move the Go-Live date

Working in the Right Environment

Environments support a variety of implementation needs

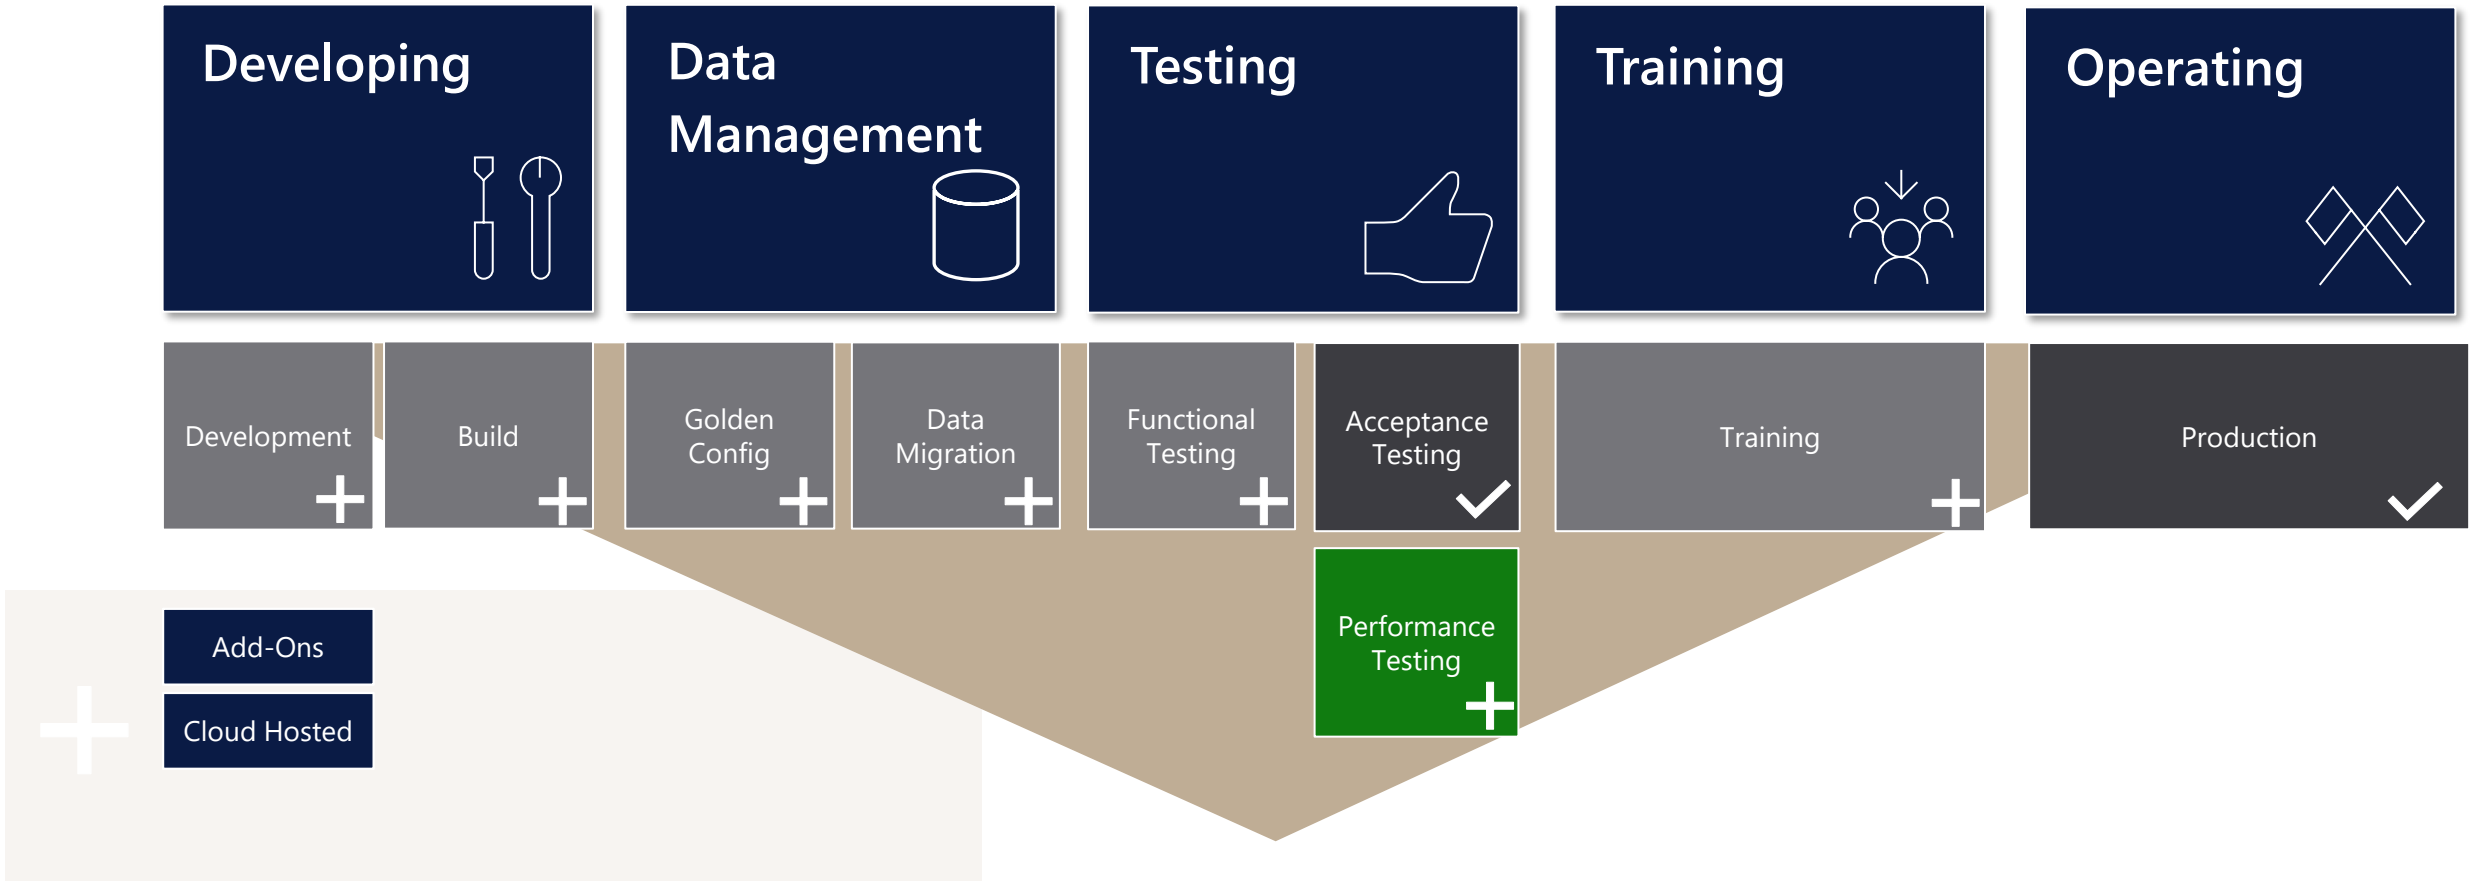

Selecting the correct Tier-2+ for performance testing

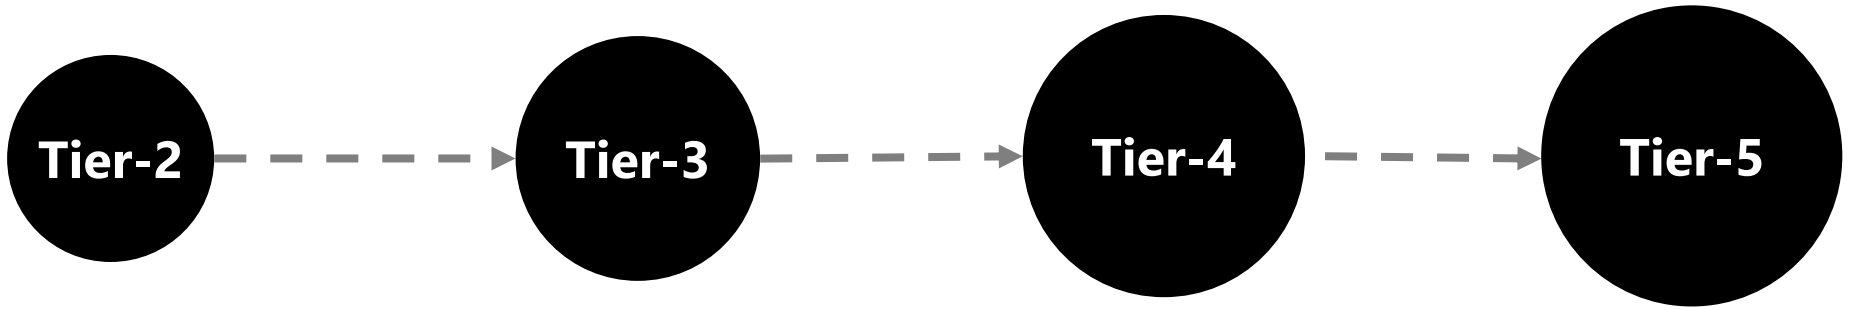

Standard Acceptance Test	Premier Acceptance Test	Standard Performance Test	Premier Performance Test
30-40 users or **3,000 transaction lines per hour****	**40-70 users** or **7,000 transaction lines per hour****	**70-150 users** or **20,000 transaction lines per hour****	**Above 150 users** or **above 20,000 transaction lines per hour****

Rely on your **partner** to adjust this **baseline** considering the specific **business scenarios** (type of users/complexity/volumes/…)

**Note: Transaction lines can be found after uploading a usage profile in the Subscription Estimator

[Environment planning - Finance & Operations | Dynamics 365 | Microsoft Docs]

Tools

Latency Test

There are handy websites like AzureSpeed.com that help with latency testing.

High latency can influence UX negatively.

Review System requirements for cloud deployments (Docs) for latency and bandwidth requirements.

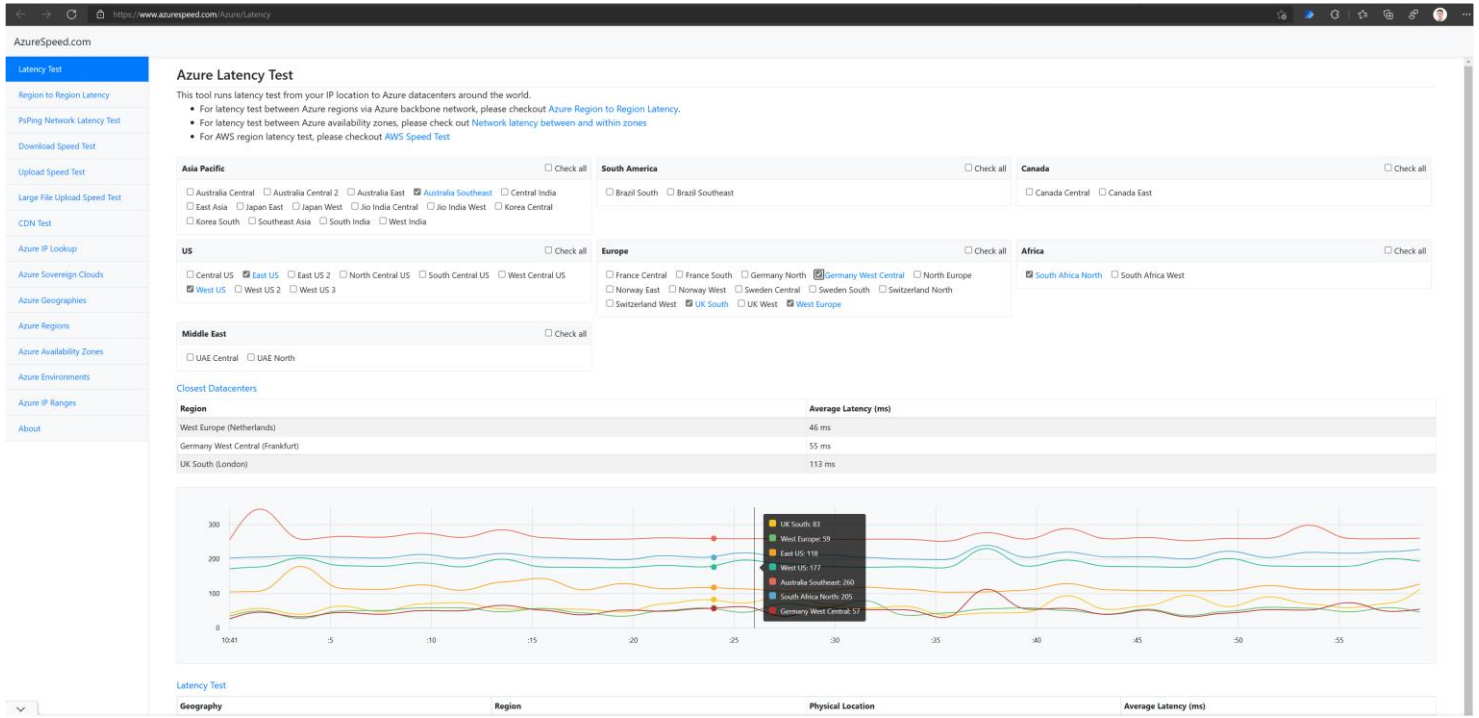

Session information

The *Feedback* dialog gives you session and activity Ids.

On Tier-2+ and Production environments, you can use these to investigate telemetry using Lifecycle Services.

These information are crucial to provide when opening support tickets with Microsoft Support too.

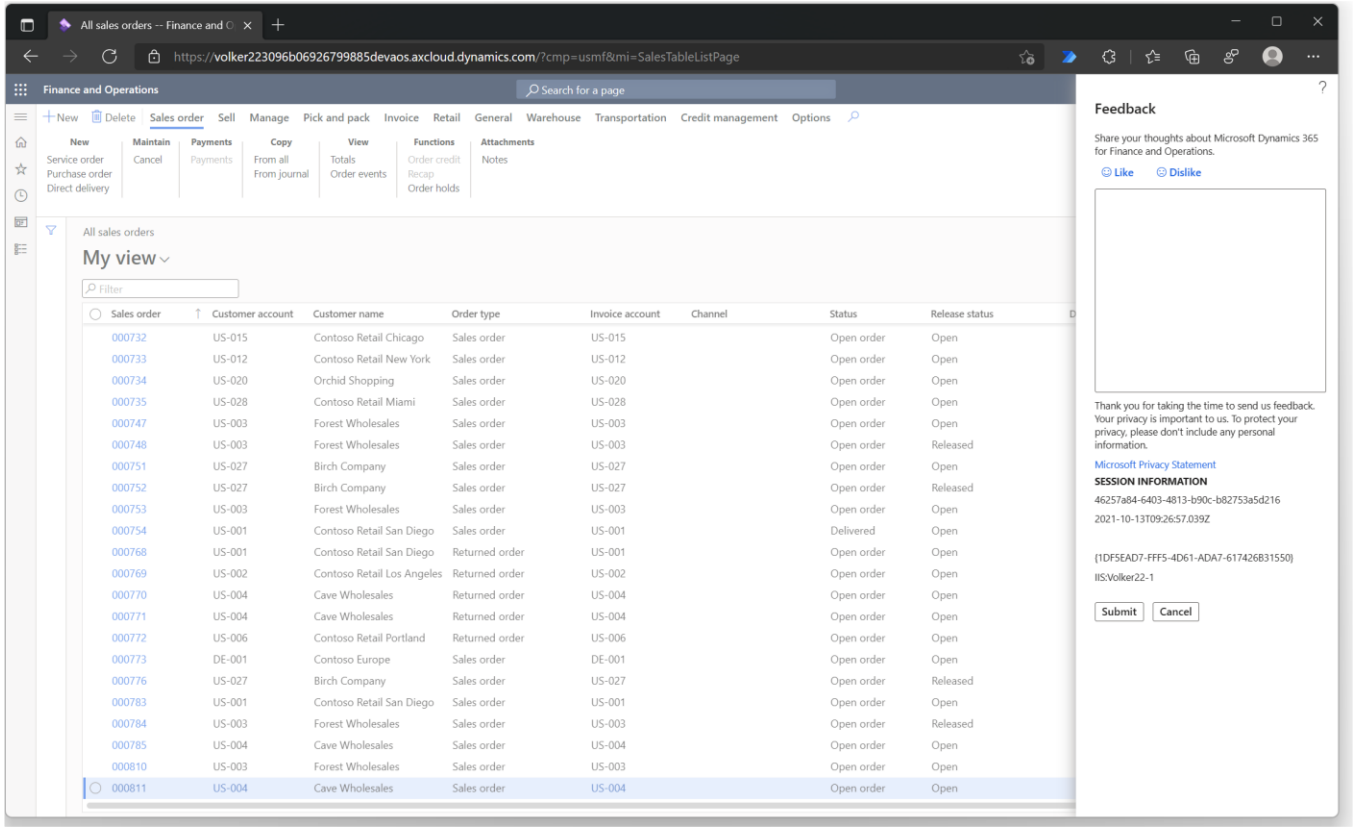

Lifecycle Services Monitoring and Diagnostics

The tools available on LCS might be more powerful than you think!

You can track sessions, activities, look at summarized information like Slow queries, and many other useful metrics.

See [Monitoring and diagnostics tools in Lifecycle Services (LCS)] on Docs for further information and start utilizing.

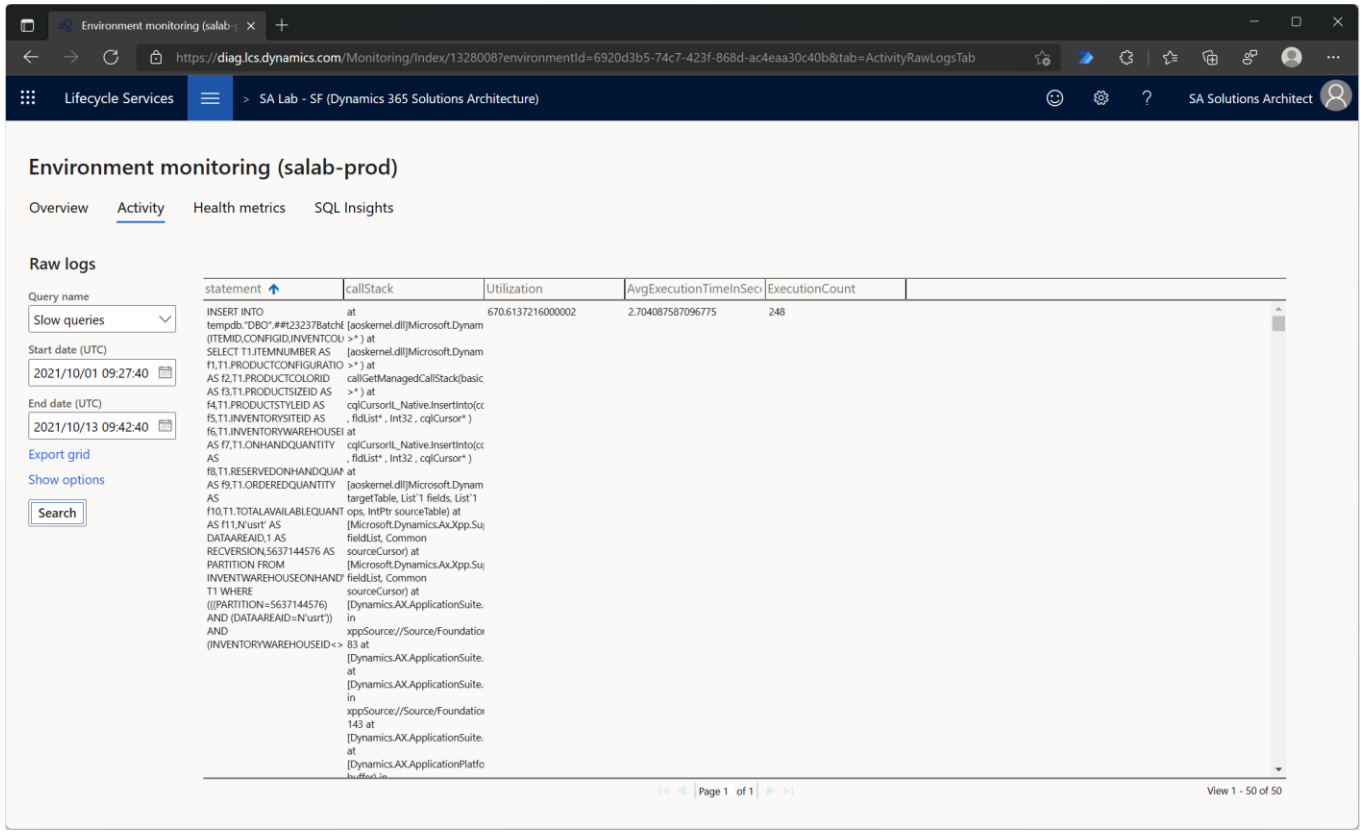

Azure Data Explorer

Freshly documented, there is a handy way to ingest the log data from LCS into an Azure Data Explorer Cluster and use KQL (Kusto Query Language) to do more complex, custom querying on the telemetry data.

See the Use Azure Data Explorer to query raw information logs tutorial on Docs.

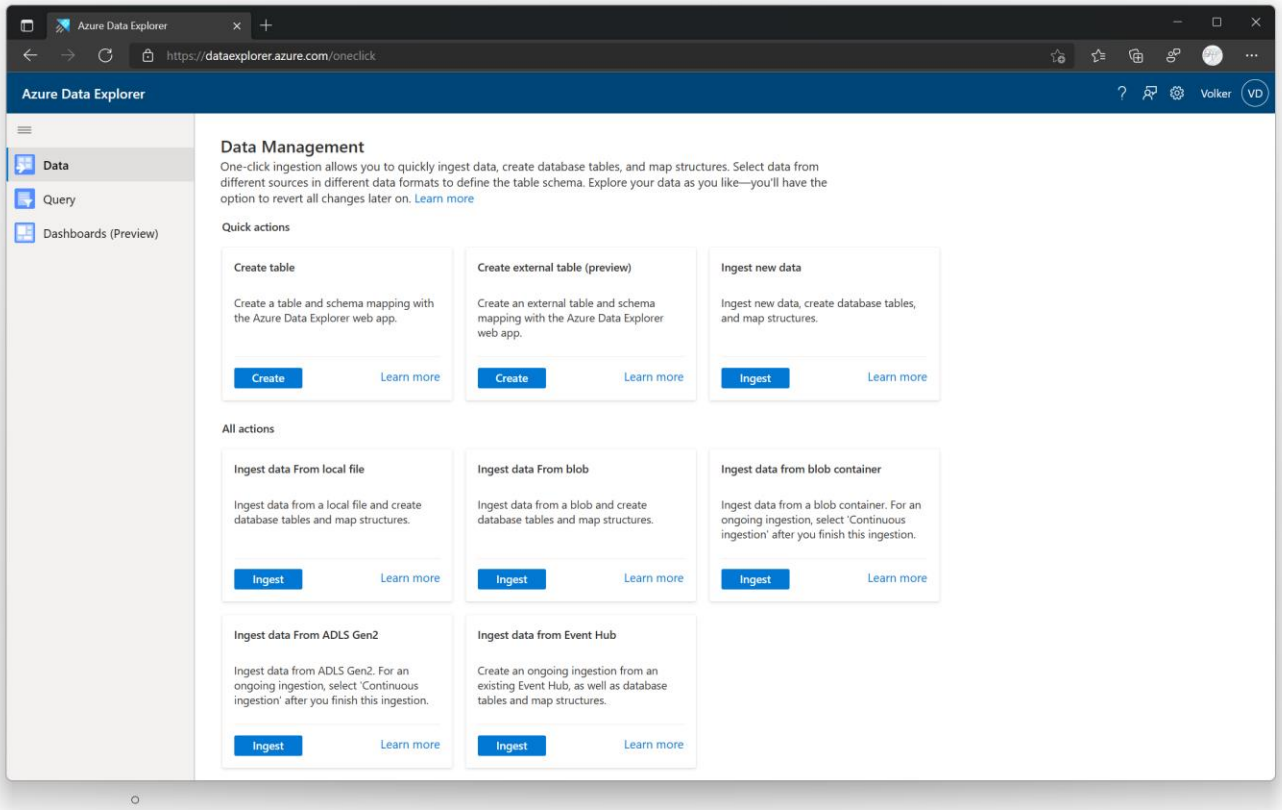

SQL Server Tools

- SSMS & Query Store [link]
- Database Engine Tuning Advisor [link]
- Profiler [link]

Query Store gets copied from Production. Use LCS to allowlist your IP and get JIT access.

Database Engine Tuning Advisor can be useful as input, but keep in mind it is not aware of the typical database schema of Finance and Operations apps.

Profiler is most useful on dev boxes.

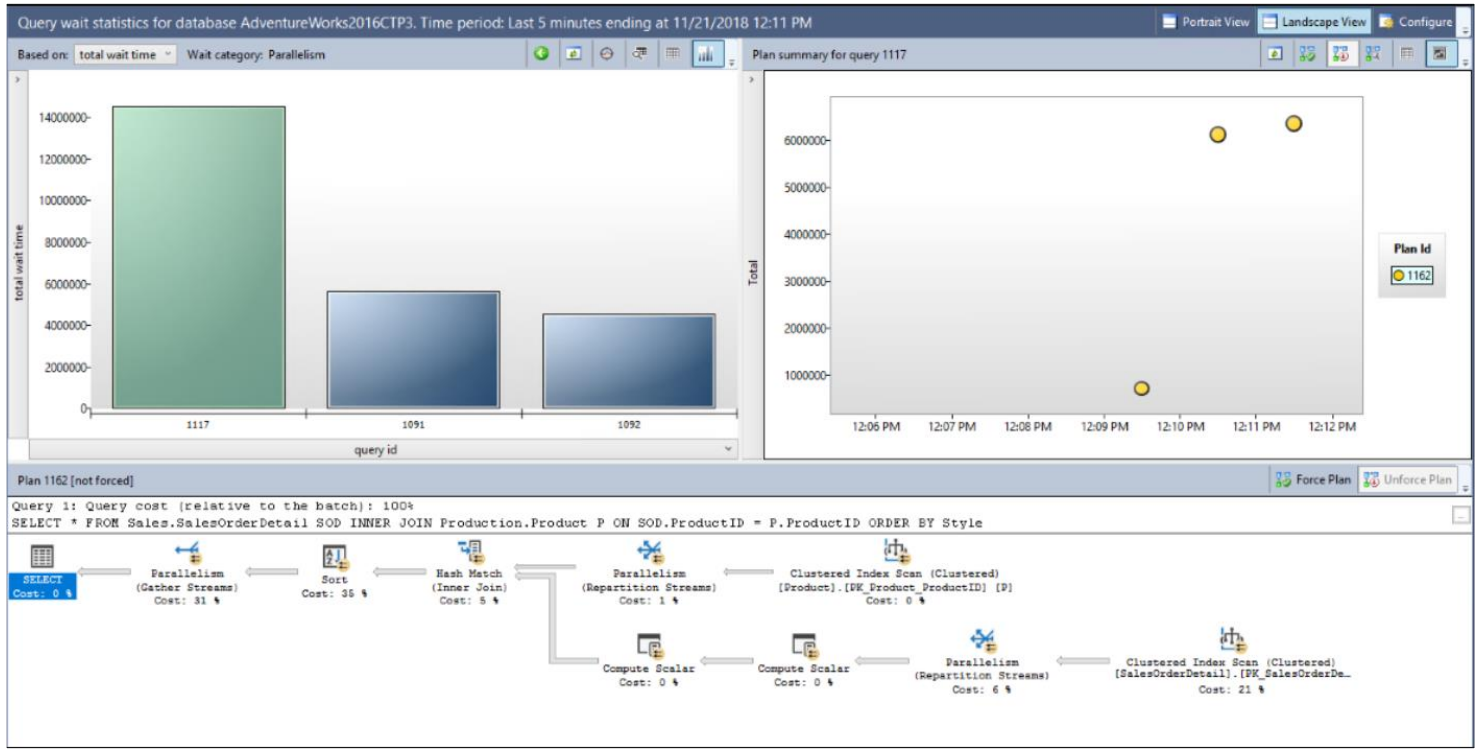

Customization Analysis Report

Many issues and warnings listed in the CAR relate to potential performance impacts.

You create it the CAR using the command line in a development environment.

See Customization Analysis Report on Docs for further information.

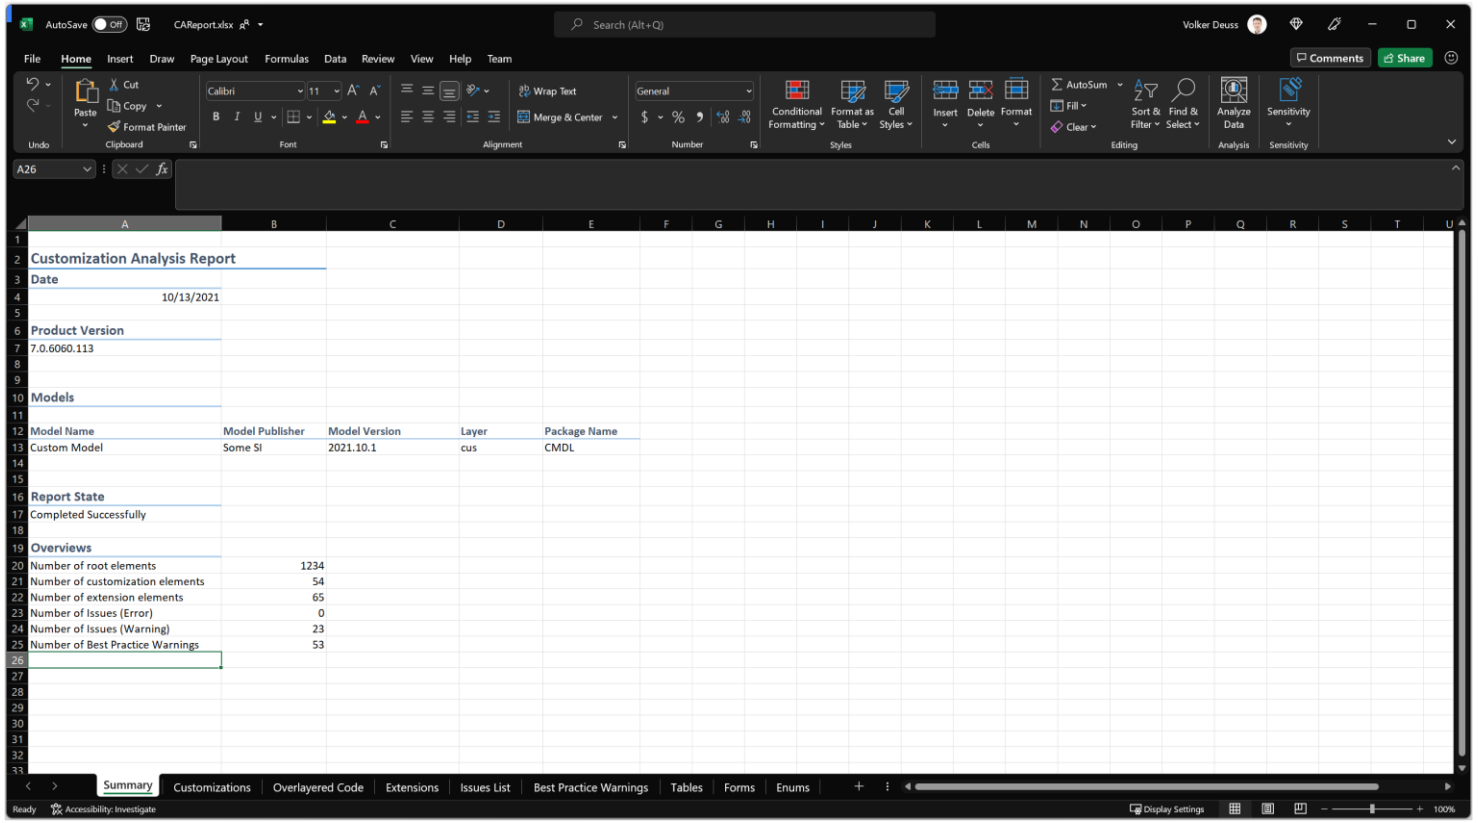

Customization Analysis Report

Date

10/13/2021

Product Version

7.0.6060.113

Models

Model Name	Model Publisher	Model Version	Layer	Package Name
Custom Model	Some SI	2021.10.1	cus	CMDL

Report State

Completed Successfully

Overviews

Number of root elements	1234
Number of customization elements	54
Number of extension elements	65
Number of Issues (Error)	0
Number of Issues (Warning)	23
Number of Best Practice Warnings	53

Load Test Automation

Performance SDK is one example and tool* to automate performance testing and benchmarking.

Load test automation helps you to validate, monitor, and optimize scalability of the solution.

Resources:

- TechTalk Performance Benchmark
- Docs Pages

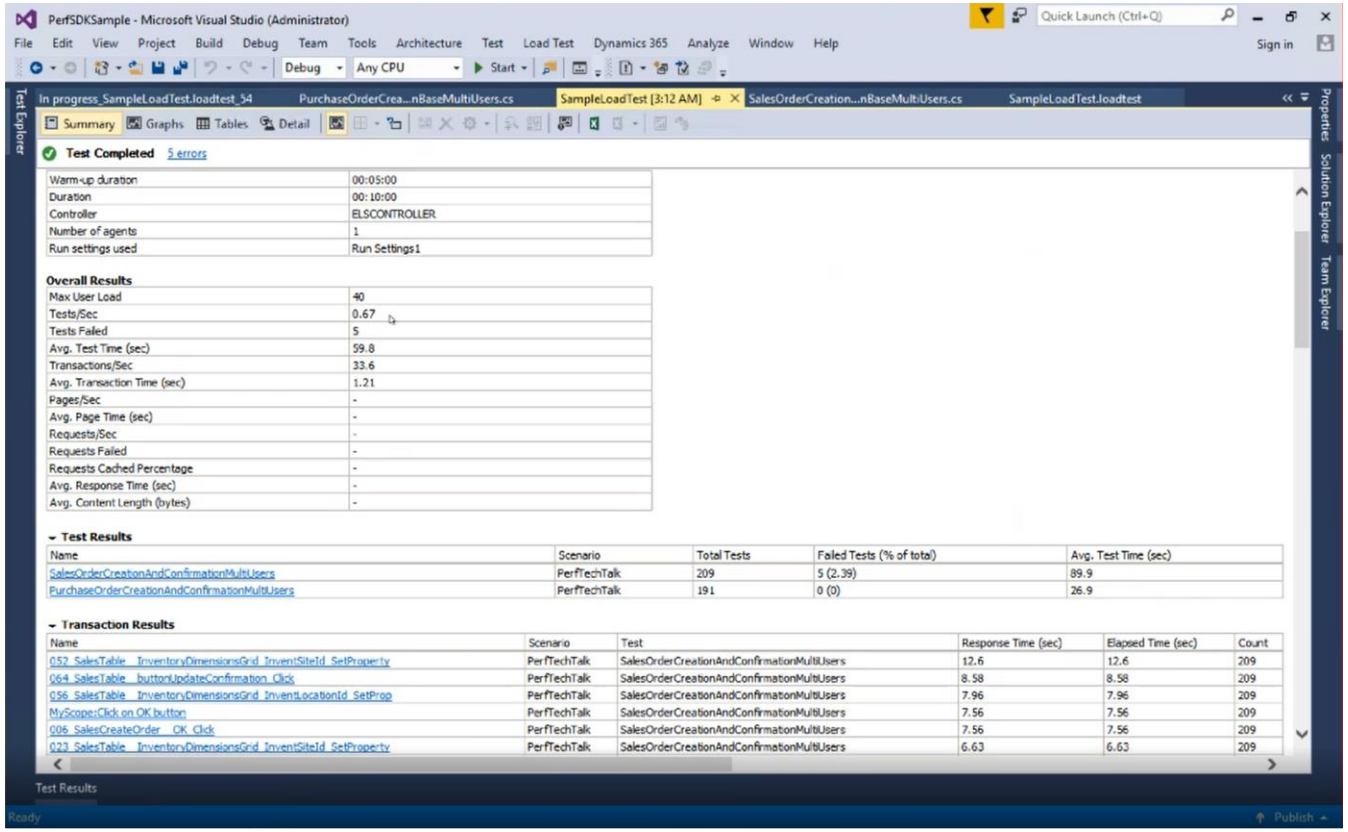

*Basically, it is just a library and set of tools on top of Visual Studio Load Test capabilities.

Performance Timer

Add &debug=develop to URL.

Built-in tooling that can help you to get an understanding of where the time is spent.

Note: Using it has negative impact on the session's performance.

See Performance timer on Docs for further information.

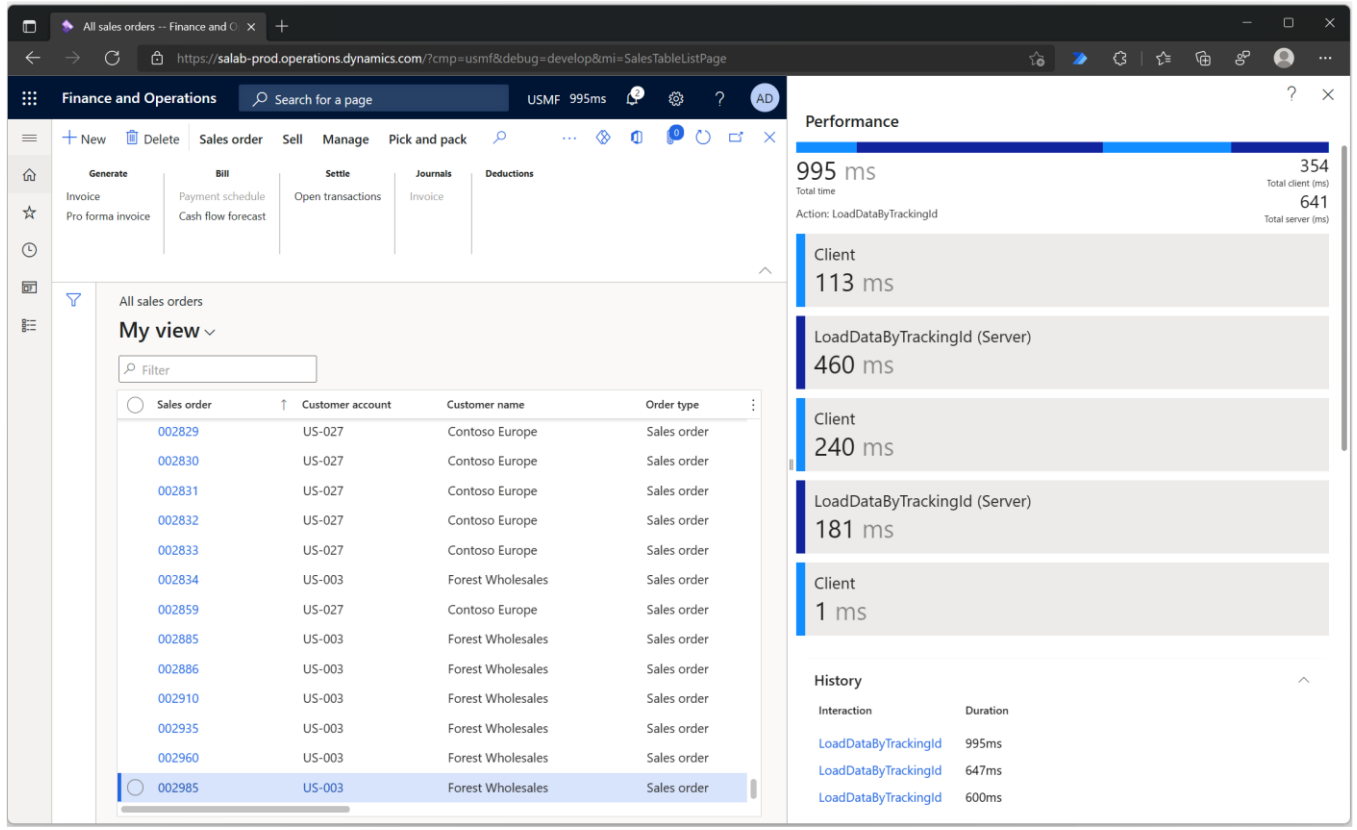

F12 Dev Tools

Modern browsers like Edge or Chrome have powerful tooling available to analyze performance aspects from the client. Typically, pressing F12 opens them.

See Microsoft Edge Developer Tools overview on Docs for Edge documentation.

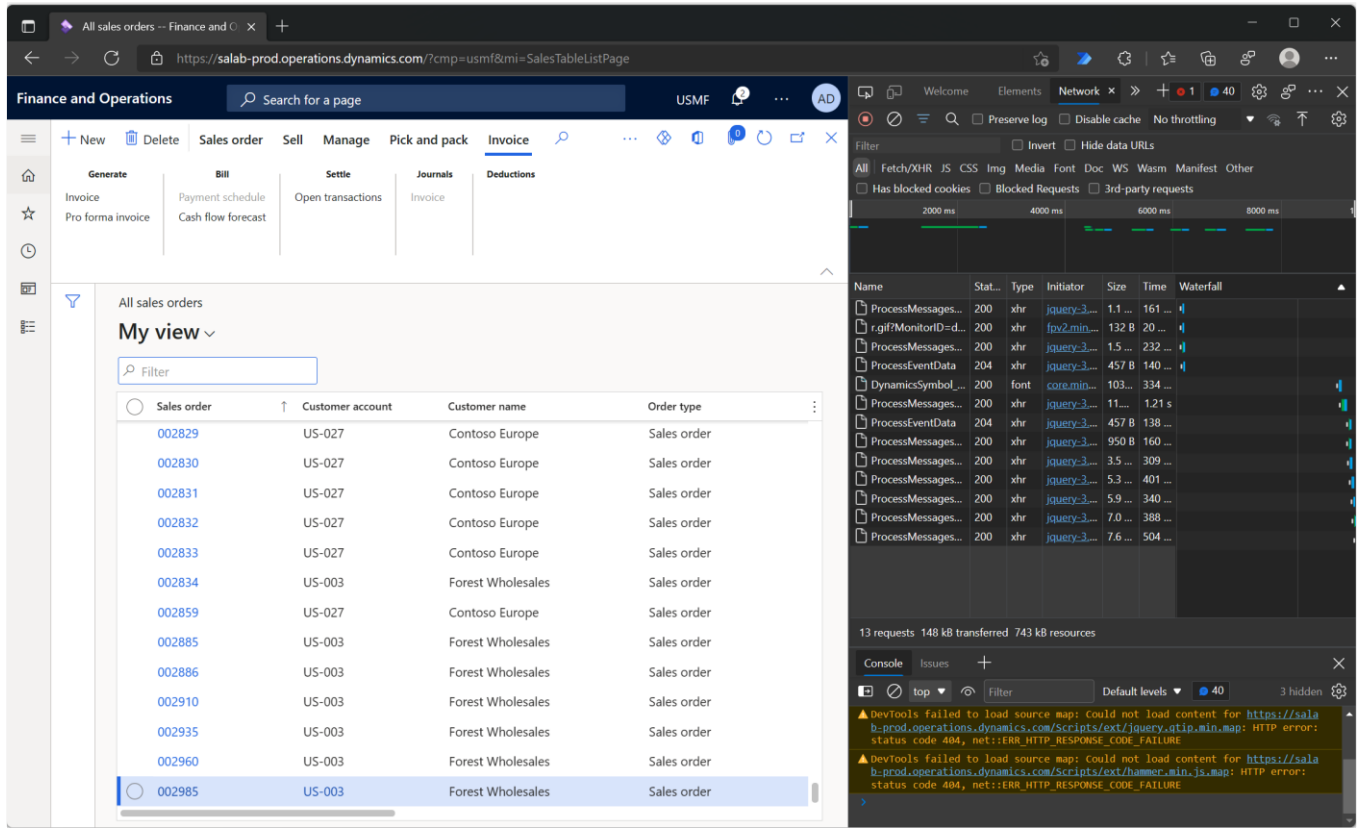

Fiddler

Powerful, widely used tool to capture and analyze traffic.

It is a 3rd party tool, available at Fiddler | Web Debugging Proxy and Troubleshooting Solutions (telerik.com).

We have a bunch of examples related to Finance and Operations apps on Docs. See Search | Microsoft Docs.

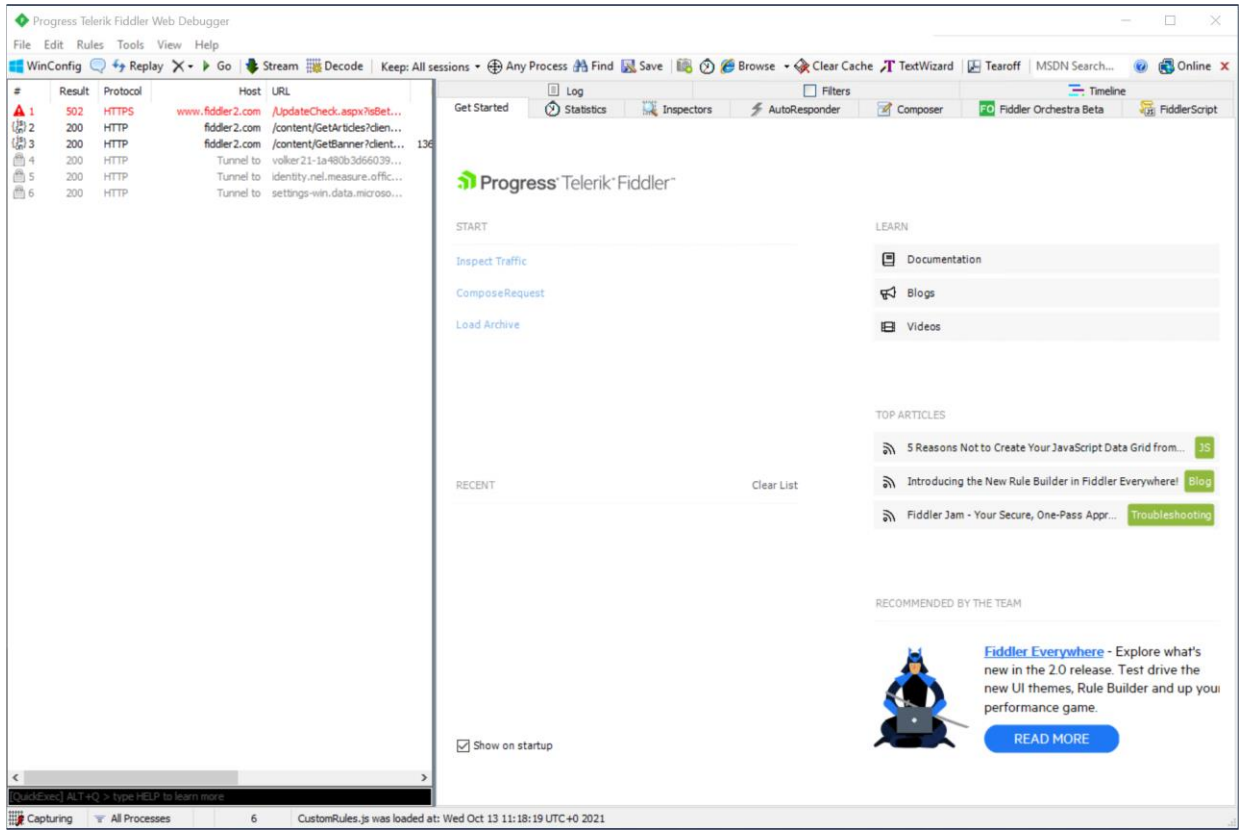

Trace Parser

Take a trace using the built-in tool available under

Help & Support / Trace

Then, use Trace Parser (installed on development machines) to analyze call trees, summarized method calls, SQL statement data, etc.

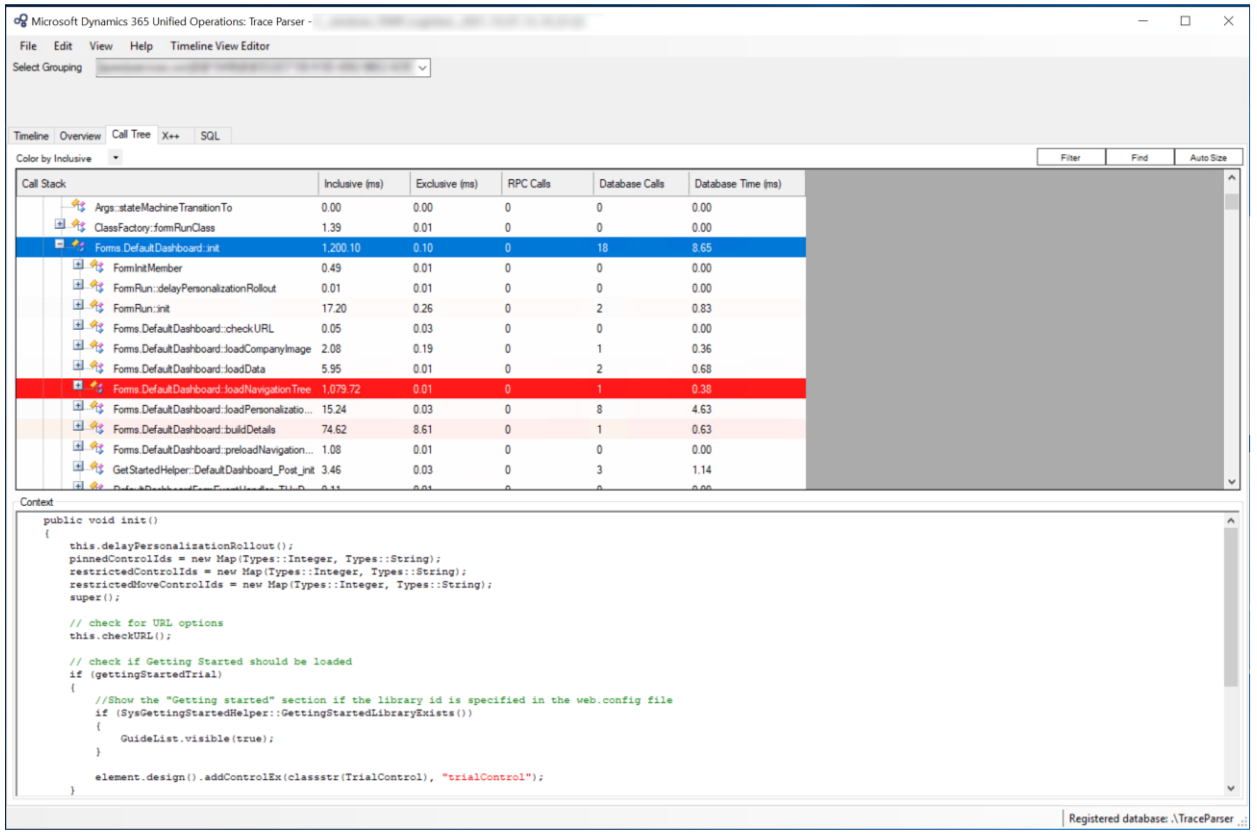

Visual Studio / Debugger

Finally, sometimes you can get to the bottom of a performance issue best by just reading the code or debugging it. It is the hard way most of the times though.

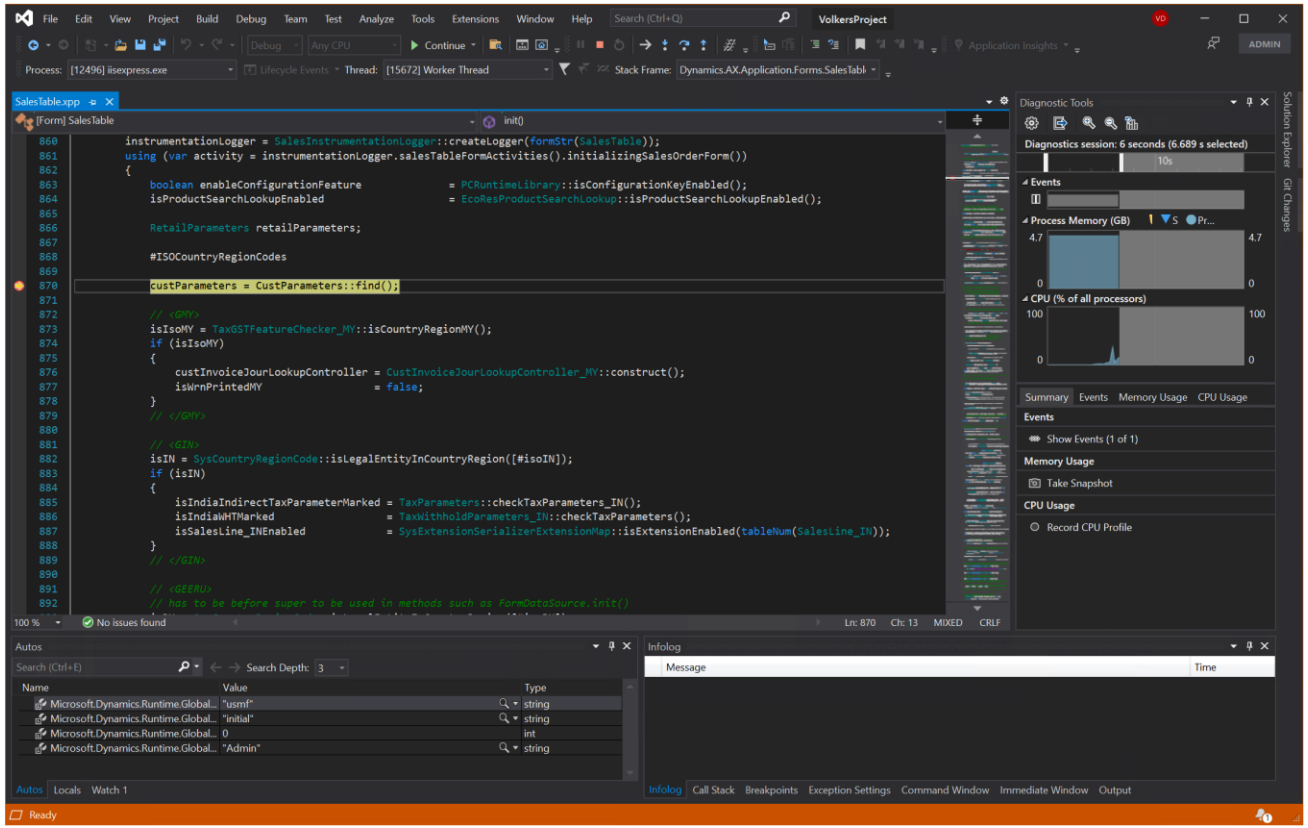

Best practices

Batch framework

In Self-Service deployment topologies, batch and interactive AOS's are isolated

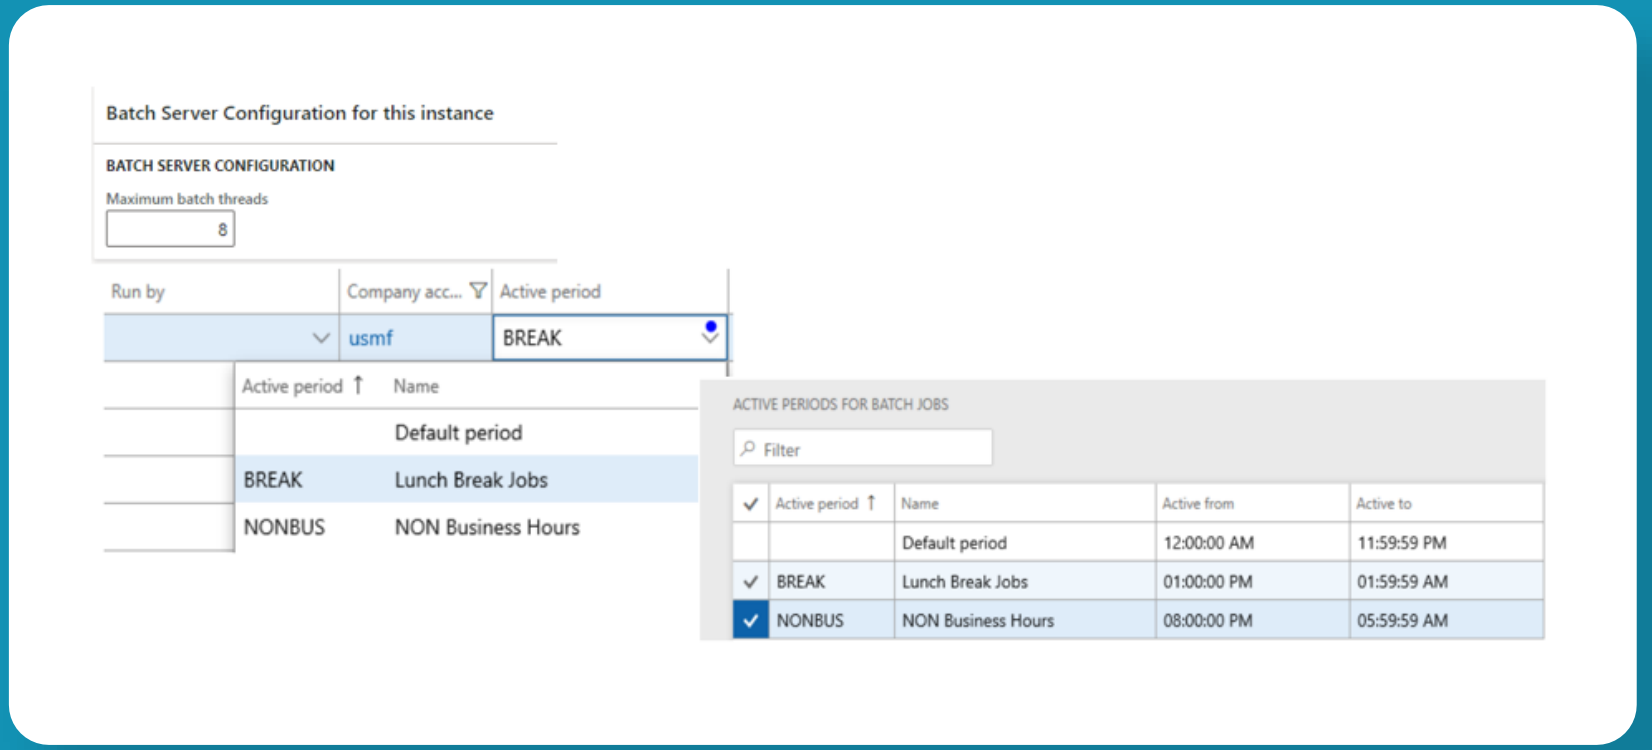

Tune your batch framework

- Create a 24-hour timetable to get an overview which heavy (batch) processes are running at a specific timeframe
- Define different active periods and assign batch jobs to decide in which time of the day the batch job can start (and when it must not)
- Do not increase the Max Batch Threads without testing (recommended value range is 8-12)
- Consider implementing automatic retries on your batch jobs
- Plan and update the scheduling priority for each batch group so that it represents the relative priorities according to business requirements for the related batch jobs and their related business using Priority based batch scheduling

Cleanup routines

- Crucial for system health and database storage
- Directly impact your system performance
- Should be run only after the business has done detailed analysis and confirmed that the data is no longer required
- Always test each cleanup routine in a test environment before you run it in a production environment

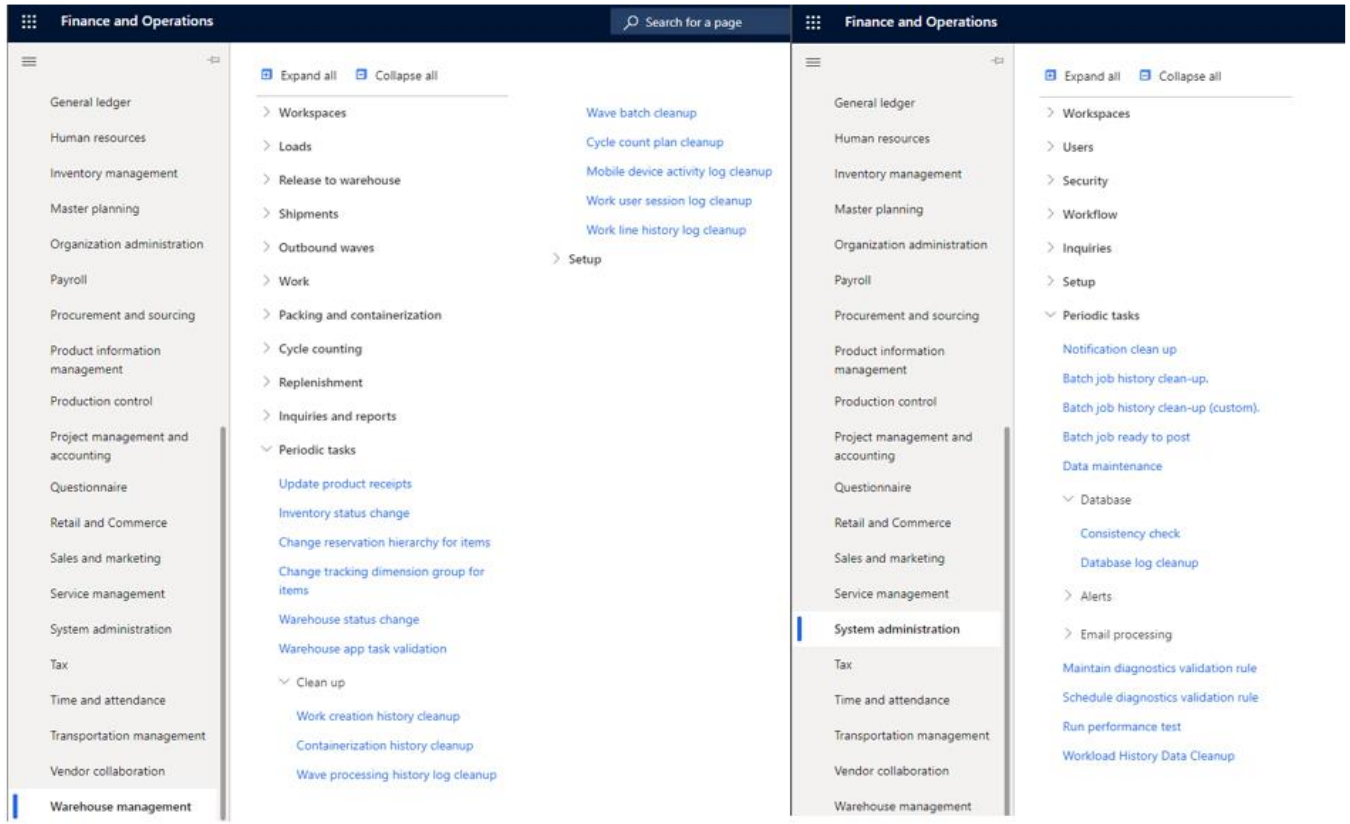

Regular cleanup routines

- Batch history tables (BatchHistory, BatchJobHistory, BatchConstraintsHistory)
  - System administration > Periodic tasks > Batch job history clean-up
- Notification tables (EventInbox, EventInboxData)
  - System administration > Periodic tasks > Notification clean up
- DMF Staging tables and Custom tables that holds temporary data
  - you must turn on the **Execution history cleanup feature** in Feature management. In Data management, this routine must be used to schedule a periodic cleanup of the execution history. It replaces the earlier Staging cleanup routine, which is now obsolete (deprecated).
  - Design a cleanup job for your temp tables
- Journal cleanup routines
  - General Ledger > Periodic tasks > Clean up ledger journals
  - Inventory management > Periodic tasks > Clean up > Inventory journals cleanup
  - Production control > Periodic tasks > Clean up > Production journals cleanup

Database Log

Database logging can be valuable from a business perspective, but it can also be expensive regarding resource use and management.

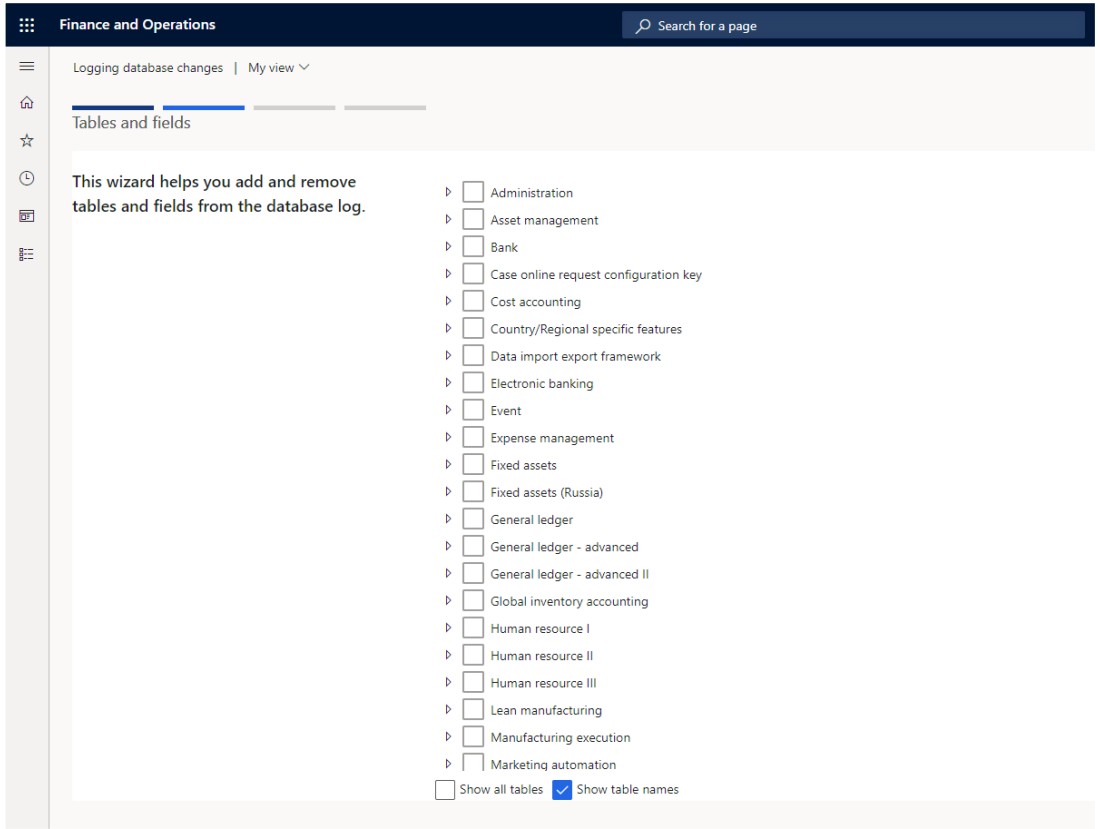

Database logging and performance

- Facts
  - The database log table can grow quickly and can increase the size of the database.
  - Database logging can adversely affect long-running automated processes, such as inventory close, calculations for bills of materials (BOMs), master planning, and long-running data imports.
  - When logging is turned on for a table, all set-based database operations are downgraded to row-based operations.
- Recommendations
  - Create a plan for how long you will retain logged data, and how you will archive or delete data
  - Limit log entries and help improve performance by selecting specific fields to log instead of whole tables,
  - Log setup data and master data and minimize logging on transactional tables.

Indexing and query performance

You can track Slow queries in LCS logs and in the SQL query store on your Tier-2 Sandbox

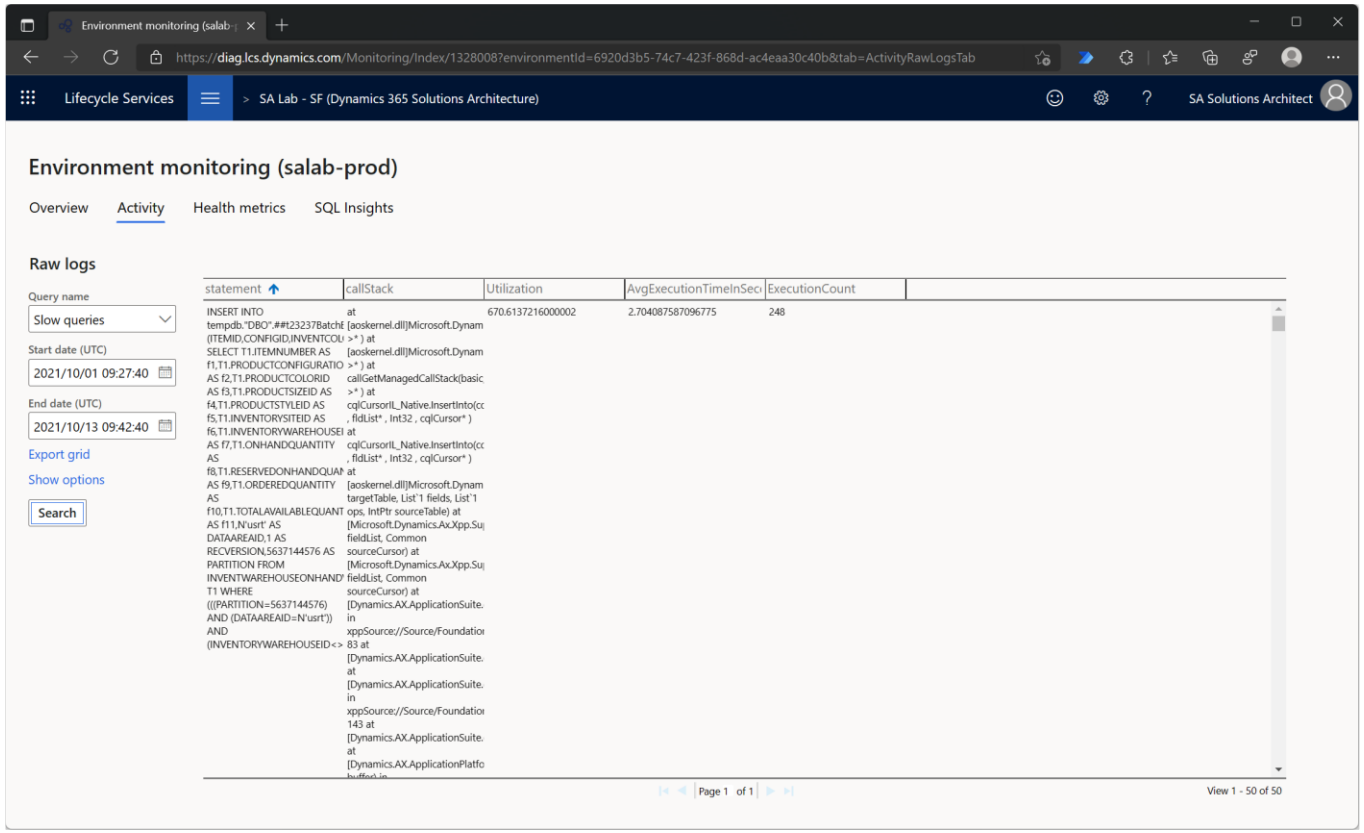

SQL performance

- Facts
  - Index scans can have negative effect on your performance and can cause blockings even by selects running in a context of a transaction, a proper index can greatly increase the query execution time .
  - Our automated microservices will take care of index maintenance jobs by rebuilding indexes when required without creating blockings and also updating statistics without the tradeoff between improving plans and the time it takes to recompile queries.
- Recommendations
  - Compare and analyze Query execution time in LCS and query store (in some cases Execution time in LCS > execution time in SQL)
  - Make sure you have the right indexes in place and that each table has a clustered index
  - Test your query execution with different data sets (different LE's…etc)

Number sequence

Consider enabling Prealloaction

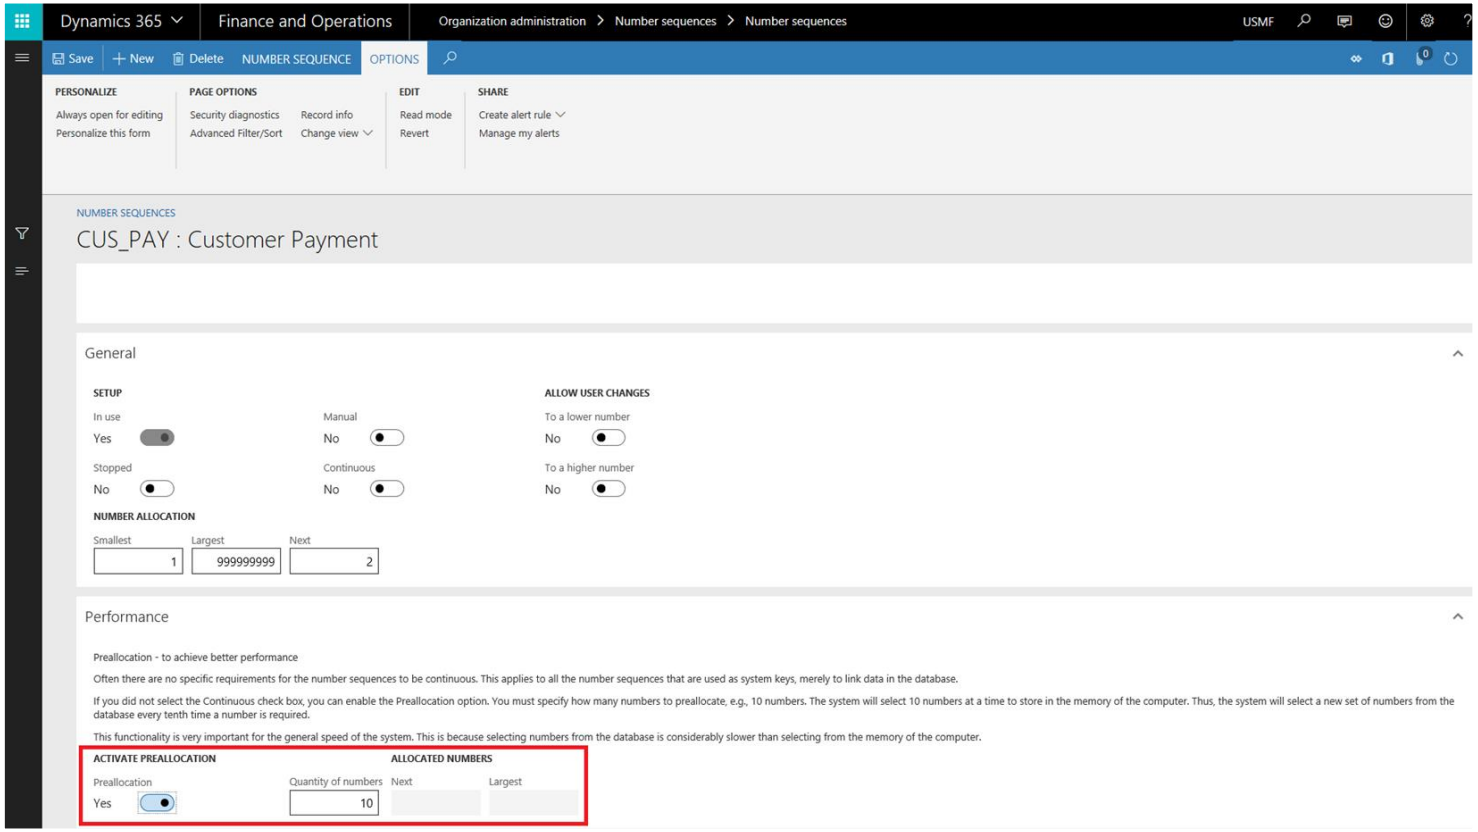

Number sequence

- Recommendations
  - Review and adjust as early in the project
  - Consider enabling Preallocation on the Performance FastTab
  - Minimize the use of the continuous number sequence as much as possible
- Pre-allocation will apply only to non-continuous Number Sequence.
  - 100 might be appropriate when 75,000+ numbers are being used each day
  - 20-50 might be appropriate when 25,000+ numbers are being used each day
  - 10 might be appropriate when 10,000+ numbers are being used each day

Data migration

Not all standard entities have been optimized for data migration.

If a required entity can't be optimized to meet the performance requirements, we recommend that you create a new optimized entity.

A developer can accelerate this process by duplicating an existing entity.

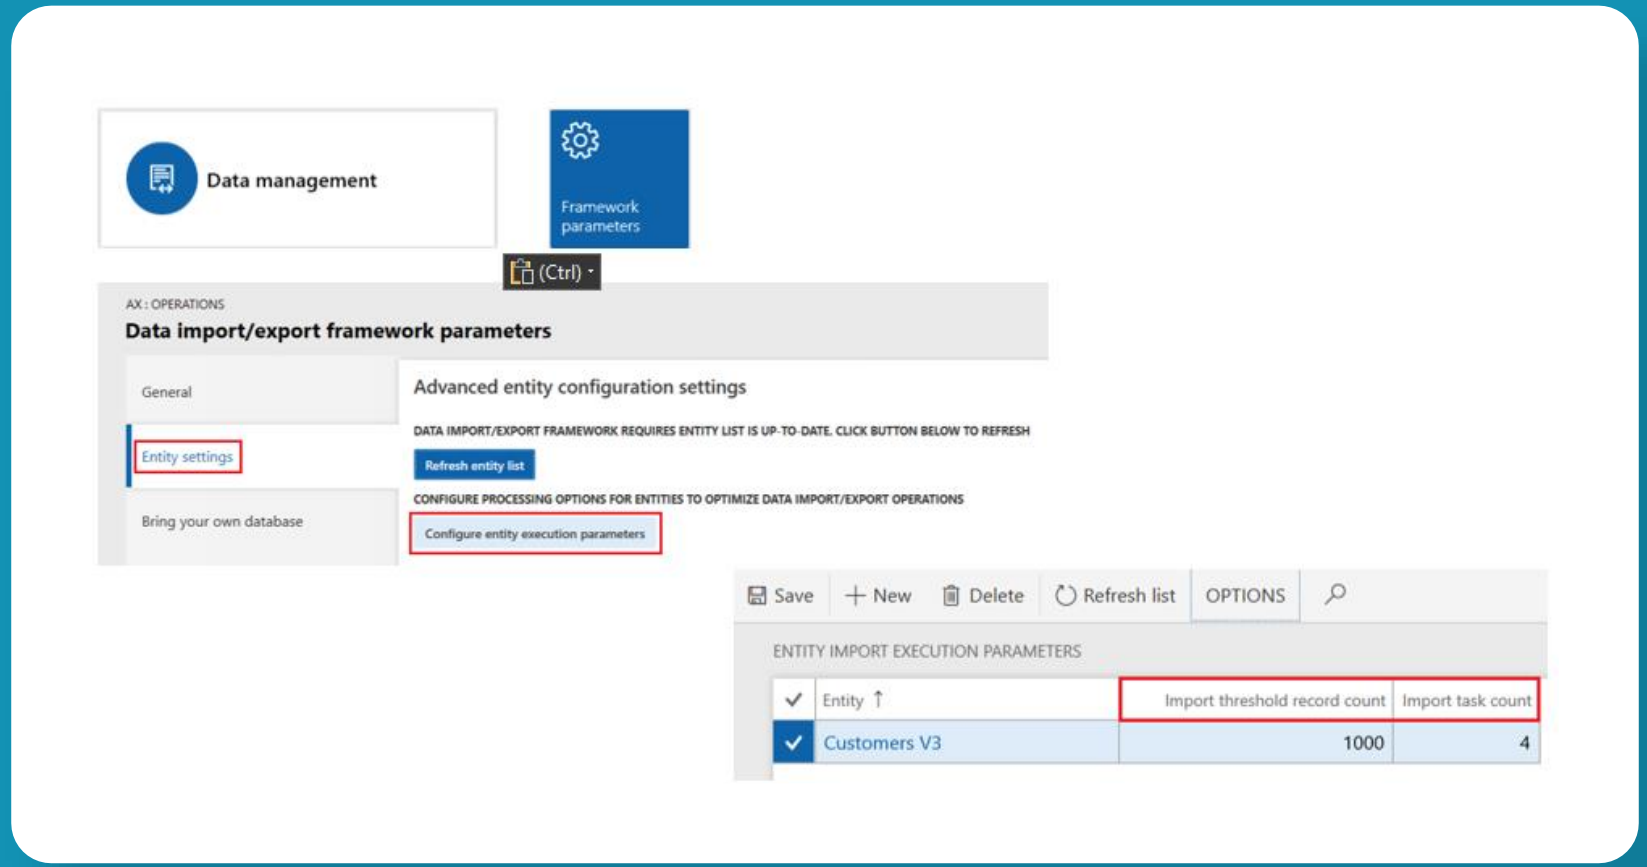

Optimize data migration

- Plan iterations of increasing size / volume
- Use delta loading and break up large files into smaller chunks. This approach gives the SQL optimizer time to determine whether a new query plan will be optimal.
- Optimize the performance parameters of the data entities for large imports; do not assume DMF entities are tuned for performance out of the box
- Clean up staging tables regularly
- Disable change tracking
- Enable set-based processing (not all entities support it)
- Create a data migration batch group and use batch mode to import to avoid single threading
- If contention is identified in the batch framework, ensure priority-based batch scheduling priorities are configured
- Update statistics
- Make sure you are using an entity that supports multithreading and configure entity execution parameters (Import record count, task count, validations as appropriate)

Resources

- Enable Batch retry Enable automatic retries on batch jobs - Finance & Operations | Dynamics 365 | Microsoft Docs
- Batch Active Periods Active batch periods - Finance & Operations | Dynamics 365 | Microsoft Docs
- Priority based batch scheduling Priority-based batch scheduling - Finance & Operations | Dynamics 365 | Microsoft Docs
- Cleanup routines Cleanup routines in Dynamics 365 Finance and Dynamics 365 Supply Chain Management - Finance & Operations | Dynamics 365 | Microsoft Docs
- Database Log Configure database logging - Finance & Operations | Dynamics 365 | Microsoft Docs
- Throttling Priority-based throttling - Finance & Operations | Dynamics 365 | Microsoft Docs
- Throttling FAQ Priority-based throttling FAQ - Finance & Operations | Dynamics 365 | Microsoft Docs
- Data Migration Optimize data migration for Finance and Operations apps - Finance & Operations | Dynamics 365 | Microsoft Docs
- Set based operations conversion Conversion of operations from set-based to record-by-record - Finance & Operations | Dynamics 365 | Microsoft Docs

- **TechTalks: Performance Testing in Microsoft Dynamics 365 TechTalk Series - Microsoft Dynamics Blog**
- See also previous sessions in this series such as Best Practices: Performance Troubleshooting | Microsoft Learn

Q&A

Microsoft Dynamics 365

© Copyright Microsoft Corporation. All rights reserved.